

Self-Driving Car Engineer Nanodegree ¶

<https://github.com/udacity/CarND-Traffic-Sign-Classifer-Project>

(<https://github.com/udacity/CarND-Traffic-Sign-Classifer-Project>)

Deep Learning

Project: Build a Traffic Sign Recognition Classifier

In this notebook, a template is provided for you to implement your functionality in stages, which is required to successfully complete this project. If additional code is required that cannot be included in the notebook, be sure that the Python code is successfully imported and included in your submission if necessary.

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the iPython Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to "\n", **"File -> Download as -> HTML (.html)"**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there is a writeup to complete. The writeup should be completed in a separate file, which can be either a markdown file or a pdf document. There is a [write up template \(https://github.com/udacity/CarND-Traffic-Sign-Classifer-Project/blob/master/writeup_template.md\)](https://github.com/udacity/CarND-Traffic-Sign-Classifer-Project/blob/master/writeup_template.md) that can be used to guide the writing process. Completing the code template and writeup template will cover all of the [rubric points \(https://review.udacity.com/#!/rubrics/481/view\)](https://review.udacity.com/#!/rubrics/481/view) for this project.

The [rubric \(https://review.udacity.com/#!/rubrics/481/view\)](https://review.udacity.com/#!/rubrics/481/view) contains "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. The stand out suggestions are optional. If you decide to pursue the "stand out suggestions", you can include the code in this iPython notebook and also discuss the results in the writeup file.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. In addition, Markdown cells can be edited by typically double-clicking the cell to enter edit mode.

Step 0: Load The Data

```
In [1]: import pickle
import numpy as np
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
# Visualizations will be shown in the notebook.
%matplotlib inline
```

```
In [2]: # Load pickled data

# TODO: Fill this in based on where you saved the training and testing

DATA_DIR = "../../data/traffic-signs/"
training_file = DATA_DIR + "train.p"
validation_file = DATA_DIR + "valid.p"
testing_file = DATA_DIR + "test.p"

with open(training_file, mode='rb') as f:
    train = pickle.load(f)
with open(validation_file, mode='rb') as f:
    valid = pickle.load(f)
with open(testing_file, mode='rb') as f:
    test = pickle.load(f)

X_train, y_train = train['features'], train['labels']
X_val, y_val = valid['features'], valid['labels']
X_test, y_test = test['features'], test['labels']
```

Step 1: Dataset Summary & Exploration

The pickled data is a dictionary with 4 key/value pairs:

- 'features' is a 4D array containing raw pixel data of the traffic sign images, (num examples, width, height, channels).
- 'labels' is a 1D array containing the label/class id of the traffic sign. The file `signnames.csv` contains id -> name mappings for each id.
- 'sizes' is a list containing tuples, (width, height) representing the original width and height the image.
- 'coords' is a list containing tuples, (x1, y1, x2, y2) representing coordinates of a bounding box around the sign in the image. **THESE COORDINATES ASSUME THE ORIGINAL IMAGE. THE PICKLED DATA CONTAINS RESIZED VERSIONS (32 by 32) OF THESE IMAGES**

Complete the basic data summary below. Use python, numpy and/or pandas methods to calculate the data summary rather than hard coding the results. For example, the [pandas shape method \(http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.shape.html\)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.shape.html) might be useful for calculating some of the summary results.

Provide a Basic Summary of the Data Set Using Python, Numpy and/or Pandas

In [3]: *# Summary of data*

```
n_train = len(X_train)
n_val = len(X_val)
n_test = len(X_test)
image_shape = X_train.shape[1:]
n_classes = len(np.unique(y_train))

print("Number of training examples =", n_train)
print("Number of examples: training: %d, validation: %d, testing: %d" %
print("Image data shape: ", image_shape)
print("Number of classes: ", n_classes)

classes, counts = np.unique(y_train, return_counts=True)
signnames_df = pd.read_csv("./signnames.csv")
signnames={int(row["ClassId"]): row["SignName"] for _, row in signnames_

print(pd.DataFrame({"class":classes, "count":counts, "name": signnames_
```

Number of training examples = 34799

Number of examples: training: 34799, validation: 4410, testing: 12630

Image data shape: (32, 32, 3)

Number of classes: 43

	class	count	name
0	0	180	Speed limit (20km/h)
1	1	1980	Speed limit (30km/h)
2	2	2010	Speed limit (50km/h)
3	3	1260	Speed limit (60km/h)
4	4	1770	Speed limit (70km/h)
5	5	1650	Speed limit (80km/h)
6	6	360	End of speed limit (80km/h)
7	7	1290	Speed limit (100km/h)
8	8	1260	Speed limit (120km/h)
9	9	1320	No passing
10	10	1800	No passing for vehicles over 3.5 metric tons
11	11	1170	Right-of-way at the next intersection
12	12	1890	Priority road
13	13	1920	Yield
14	14	690	Stop
15	15	540	No vehicles
16	16	360	Vehicles over 3.5 metric tons prohibited
17	17	990	No entry
18	18	1080	General caution
19	19	180	Dangerous curve to the left
20	20	300	Dangerous curve to the right
21	21	270	Double curve
22	22	330	Bumpy road
23	23	450	Slippery road
24	24	240	Road narrows on the right
25	25	1350	Road work
26	26	540	Traffic signals
27	27	210	Pedestrians
28	28	480	Children crossing
29	29	240	Bicycles crossing
30	30	390	Beware of ice/snow
31	31	690	Wild animals crossing
32	32	210	End of all speed and passing limits

33	33	599	Turn right ahead
34	34	360	Turn left ahead
35	35	1080	Ahead only
36	36	330	Go straight or right
37	37	180	Go straight or left
38	38	1860	Keep right
39	39	270	Keep left
40	40	300	Roundabout mandatory
41	41	210	End of no passing
42	42	210	End of no passing by vehicles over 3.5 metric ...

Include an exploratory visualization of the dataset

Visualize the German Traffic Signs Dataset using the pickled file(s). This is open ended, suggestions include: plotting traffic sign images, plotting the count of each sign, etc.

The [Matplotlib](http://matplotlib.org/) (<http://matplotlib.org/>) [examples](http://matplotlib.org/examples/index.html) (<http://matplotlib.org/examples/index.html>) and [gallery](http://matplotlib.org/gallery.html) (<http://matplotlib.org/gallery.html>) pages are a great resource for doing visualizations in Python.

NOTE: It's recommended you start with something simple first. If you wish to do more, come back to it after you've completed the rest of the sections. It can be interesting to look at the distribution of classes in the training, validation and test set. Is the distribution the same? Are there more examples of some classes than others?

```
In [4]: # Plot example traffic signs
cols = 3
rows = len(classes)
fig, axes = plt.subplots(rows, cols, figsize=(4*cols, 4*rows))
np.random.seed(1)
for i, class_id in enumerate(classes[:rows]):
    indices = np.where(y_train==class_id)[0]
    indices = np.random.choice(indices, cols, replace=False)
    for j, idx in enumerate(indices):
        axes[i, j].imshow(X_train[idx,:,:,])
        axes[i, j].set_xticks([])
        axes[i, j].set_yticks([])
    axes[i, 0].set_title("%d: %s" % (class_id, signnames[class_id]))
```



Step 2: Design and Test a Model Architecture

Design and implement a deep learning model that learns to recognize traffic signs. Train and test your model on the [German Traffic Sign Dataset](http://benchmark.ini.rub.de/?section=gtsrb&subsection=dataset) (<http://benchmark.ini.rub.de/?section=gtsrb&subsection=dataset>).

The LeNet-5 implementation shown in the [classroom](https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/601ae704-1035-4287-8b11-e2c2716217ad/concepts/d4aca031-508f-4e0b-b493-e7b706120f81) (<https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/601ae704-1035-4287-8b11-e2c2716217ad/concepts/d4aca031-508f-4e0b-b493-e7b706120f81>) at the end of the CNN lesson is a solid starting point. You'll have to change the number of classes and possibly the preprocessing, but aside from that it's plug and play!

With the LeNet-5 solution from the lecture, you should expect a validation set accuracy of about 0.89. To meet specifications, the validation set accuracy will need to be at least 0.93. It is possible to get an even higher accuracy, but 0.93 is the minimum for a successful project

submission.

There are various aspects to consider when thinking about this problem:

- Neural network architecture (is the network over or underfitting?)
- Play around preprocessing techniques (normalization, rgb to grayscale, etc)
- Number of examples per label (some have more than others).
- Generate fake data.

Here is an example of a [published baseline model on this problem](http://yann.lecun.com/exdb/publis/pdf/sermanet-ijcnn-11.pdf)

(<http://yann.lecun.com/exdb/publis/pdf/sermanet-ijcnn-11.pdf>). It's not required to be familiar with the approach used in the paper but, it's good practice to try to read papers like these.

Pre-process the Data Set (normalization, grayscale, etc.)

Minimally, the image data should be normalized so that the data has mean zero and equal variance. For image data, $(\text{pixel} - 128) / 128$ is a quick way to approximately normalize the data and can be used in this project.

Other pre-processing steps are optional. You can try different techniques to see if it improves performance.

Use the code cell (or multiple code cells, if necessary) to implement the first step of your project.

```
In [5]: # Preprocess: normalize, grayscale?
def normalize(X):
    return (X - 128.0) / 128.0

X_train = normalize(X_train)
X_val = normalize(X_val)
X_test = normalize(X_test)
#print(X_test.mean())
```

Model Architecture

```

In [6]: # Model
import tensorflow as tf

def TrafficSignModel(X, n_classes, dropout, is_training):
    # Conv1: input: 32x32x3, output: 28x28x6
    conv1 = conv2d(X, ksize=5, stride=1, in_channels=3, out_channels=6)
    conv1 = tf.nn.relu(conv1)
    conv1 = max_pool(conv1, ksize=2, stride=2)
    conv1 = tf.layers.dropout(conv1, rate=dropout, training=is_training)

    # Conv2: input: 14x14x6, output: 10x10x16
    conv2 = conv2d(conv1, ksize=5, stride=1, in_channels=6, out_channels=16)
    conv2 = tf.nn.relu(conv2)
    conv2 = max_pool(conv2, ksize=2, stride=2)
    conv2 = tf.layers.dropout(conv2, rate=dropout, training=is_training)

    # Input: 5x5x16, output: 400
    fc0 = tf.layers.flatten(conv2)

    # FC1: input: 400, output: 120
    fc1_W = tf.Variable(tf.truncated_normal(shape=(400, 120), mean=0, stddev=0.01))
    fc1_b = tf.Variable(tf.zeros(120))
    fc1 = tf.matmul(fc0, fc1_W) + fc1_b
    fc1 = tf.nn.relu(fc1)
    fc1 = tf.layers.dropout(fc1, rate=dropout, training=is_training)

    # FC2:
    fc2_W = tf.Variable(tf.truncated_normal(shape=(120, 84), mean=0, stddev=0.01))
    fc2_b = tf.Variable(tf.zeros(84))
    fc2 = tf.matmul(fc1, fc2_W) + fc2_b
    fc2 = tf.nn.relu(fc2)

    fc3_W = tf.Variable(tf.truncated_normal(shape=(84, n_classes), mean=0, stddev=0.01))
    fc3_b = tf.Variable(tf.zeros(n_classes))
    logits = tf.matmul(fc2, fc3_W) + fc3_b

    return logits

def conv2d(input, ksize, stride, in_channels, out_channels, padding="VALID"):
    W = tf.Variable(tf.truncated_normal(shape=(ksize, ksize, in_channels, out_channels), mean=0, stddev=0.01))
    b = tf.Variable(tf.zeros(out_channels))
    conv = tf.nn.conv2d(input, W, strides=[1, stride, stride, 1], padding=padding) + b
    return conv

def max_pool(input, ksize, stride, padding="VALID"):
    return tf.nn.max_pool(input, ksize=[1, ksize, ksize, 1], strides=[1, stride, stride, 1], padding=padding)

```

Train, Validate and Test the Model

A validation set can be used to assess how well the model is performing. A low accuracy on the training and validation sets imply underfitting. A high accuracy on the training set but low accuracy on the validation set implies overfitting.


```

In [7]: ### Train your model here.
### Calculate and report the accuracy on the training and validation set
### Once a final model architecture is selected,
### the accuracy on the test set should be calculated and reported as well
### Feel free to use as many code cells as needed.
from sklearn.utils import shuffle
import matplotlib.pyplot as plt

def train(n_classes, X_train, y_train, X_val, y_val,
         epochs=1, learning_rate=0.001, batch_size=64, dropout=0.0, seed=1234):
    tf.set_random_seed(seed)
    X = tf.placeholder(tf.float32, (None, 32, 32, 3))
    y = tf.placeholder(tf.int32, (None))
    y_one_hot = tf.one_hot(y, n_classes)
    is_training = tf.placeholder_with_default(False, shape=())

    logits = TrafficSignModel(X, n_classes, dropout, is_training)
    cross_entropy = tf.nn.softmax_cross_entropy_with_logits(labels=y_one_hot, logits=logits)
    cost = tf.reduce_mean(cross_entropy)
    optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
    prediction_op = tf.cast(tf.argmax(logits, 1), tf.int32)
    correct_op = tf.equal(prediction_op, y)
    accuracy_op = tf.reduce_mean(tf.cast(correct_op, tf.float32))
    saver = tf.train.Saver()

    metrics={"train_loss": [], "val_loss": [], "cost": []}
    with tf.Session() as sess:
        sess.run(tf.global_variables_initializer())
        n_examples = len(X_train)
        for i in range(epochs):
            X_train, y_train = shuffle(X_train, y_train)
            for offset in range(0, n_examples, batch_size):
                end = offset + batch_size
                batch_X, batch_y = X_train[offset:end], y_train[offset:end]
                sess.run(optimizer, feed_dict={X: batch_X, y: batch_y, is_training: True})
            train_accuracy = accuracy_op.eval({X: X_train, y: y_train})
            validation_accuracy = accuracy_op.eval({X: X_val, y: y_val})
            metrics["train_loss"].append(1.0 - train_accuracy)
            metrics["val_loss"].append(1.0 - validation_accuracy)
            print("Epoch %2d: train accuracy: %.3f, validation accuracy: %.3f" % (i,
                train_accuracy, validation_accuracy))
            saver.save(sess, DATA_DIR + "traffic_signs_net")
            print("Model saved.")
    vars = {
        "X": X,
        "y": y,
        "logits": logits,
        "prediction_op": prediction_op,
        "accuracy_op": accuracy_op
    }
    return vars, metrics

def evaluate(vars, X_data, y_data, batch_size=64):
    n_examples = len(X_data)
    total_accuracy = 0.0
    sess = tf.get_default_session()

```

```

for offset in range(0, n_examples, batch_size):
    batch_X, batch_y = X_data[offset : offset + batch_size], y_data[offset : offset + batch_size]
    accuracy = sess.run(vars["accuracy_op"], feed_dict={vars["X"]: batch_X, vars["y"]: batch_y})
    total_accuracy += accuracy * len(batch_X)
total_accuracy /= n_examples
return total_accuracy

epochs = 20
vars, metrics = train(n_classes, X_train, y_train, X_val, y_val, epochs)

plt.plot(range(epochs), metrics["train_loss"], label="train")
plt.plot(range(epochs), metrics["val_loss"], label="validation")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.show()

def test_model(vars, X_test, y_test):
    saver = tf.train.Saver()
    with tf.Session() as sess:
        saver.restore(sess, DATA_DIR + "traffic_signs_net")
        test_accuracy = evaluate(vars, X_test, y_test)
        print("Test accuracy: %.3f\n" % test_accuracy)

test_model(vars, X_test, y_test)

```

WARNING:tensorflow:From <ipython-input-7-f2f5b73a22b4>:18: softmax_cross_entropy_with_logits (from tensorflow.python.ops.nn_ops) is deprecated and will be removed in a future version.
Instructions for updating:

Future major versions of TensorFlow will allow gradients to flow into the labels input on backprop by default.

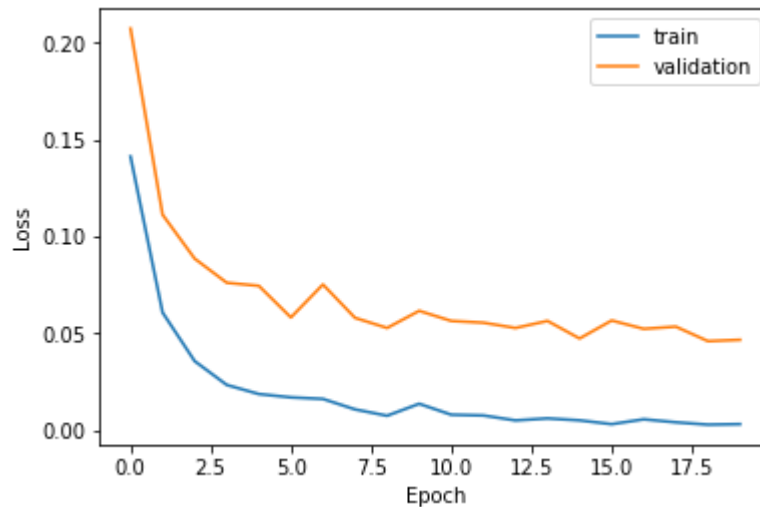
See ``tf.nn.softmax_cross_entropy_with_logits_v2``.

```

Epoch 0: train accuracy: 0.859, validation accuracy: 0.793
Epoch 1: train accuracy: 0.939, validation accuracy: 0.889
Epoch 2: train accuracy: 0.964, validation accuracy: 0.911
Epoch 3: train accuracy: 0.977, validation accuracy: 0.924
Epoch 4: train accuracy: 0.981, validation accuracy: 0.925
Epoch 5: train accuracy: 0.983, validation accuracy: 0.942
Epoch 6: train accuracy: 0.984, validation accuracy: 0.925
Epoch 7: train accuracy: 0.989, validation accuracy: 0.942
Epoch 8: train accuracy: 0.992, validation accuracy: 0.947
Epoch 9: train accuracy: 0.986, validation accuracy: 0.938
Epoch 10: train accuracy: 0.992, validation accuracy: 0.944
Epoch 11: train accuracy: 0.992, validation accuracy: 0.944
Epoch 12: train accuracy: 0.995, validation accuracy: 0.947
Epoch 13: train accuracy: 0.994, validation accuracy: 0.944
Epoch 14: train accuracy: 0.995, validation accuracy: 0.953
Epoch 15: train accuracy: 0.997, validation accuracy: 0.943
Epoch 16: train accuracy: 0.994, validation accuracy: 0.948
Epoch 17: train accuracy: 0.996, validation accuracy: 0.946
Epoch 18: train accuracy: 0.997, validation accuracy: 0.954

```

Epoch 19: train accuracy: 0.997, validation accuracy: 0.953
Model saved.



INFO:tensorflow:Restoring parameters from ../../data/traffic-signs/traffic_signs_net
Test accuracy: 0.940

Step 3: Test a Model on New Images

To give yourself more insight into how your model is working, download at least five pictures of German traffic signs from the web and use your model to predict the traffic sign type.

You may find `signnames.csv` useful as it contains mappings from the class id (integer) to the actual sign name.

Load and Output the Images

```
In [8]: ### Load the images and plot them here.  
### Feel free to use as many code cells as needed.  
import glob  
import cv2  
import os  
  
image_files = glob.glob( "data/*")  
image_files.sort()  
rows = len(image_files)  
fig, axes = plt.subplots(rows, 1, figsize=(2, 2*rows), squeeze=False)  
images = []  
for i, image_file in enumerate(image_files):  
    img = cv2.imread(image_file)  
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)  
    axes[i, 0].imshow(img)  
    img_name = os.path.basename(image_file)  
    images.append([img, img_name])  
    axes[i, 0].set_title(img_name)  
    axes[i, 0].set_xticks([])  
    axes[i, 0].set_yticks([])
```

no-entry.jpg



pedestrians.jpg



speed-limit-30.jpeg



stop_sign.jpg



warning.png



yield.jpg



Predict the Sign Type for Each Image

```

In [9]: ### Run the predictions here and use the model to output the prediction
### Make sure to pre-process the images with the same pre-processing pipeline
### Feel free to use as many code cells as needed.
def predict_images(vars, images):
    saver = tf.train.Saver()
    logits = []
    predictions = []
    with tf.Session() as sess:
        saver.restore(sess, DATA_DIR + "traffic_signs_net")
        rows = len(images)
        fig, axes = plt.subplots(rows, 1, figsize=(2, 2*rows), squeeze=False)
        for i, (raw_img, name) in enumerate(images):
            img = cv2.resize(raw_img, (32, 32))
            img = normalize(img)
            logit, prediction = sess.run([vars["logits"], vars["prediction_operation"]],
                                         feed_dict={vars["X"]: img.reshape(1, 32, 32, 3)})
            logits.append(logit[0])
            predictions.append(prediction[0])
            axes[i, 0].imshow(raw_img)
            title = "%s: class=%d, %s" % (name, prediction, signnames[prediction])
            axes[i, 0].set_title(title)
            axes[i, 0].set_xticks([])
            axes[i, 0].set_yticks([])
    return np.array(logits), np.array(predictions)

logits, predictions = predict_images(vars, images)

```

```

INFO:tensorflow:Restoring parameters from ../../data/traffic-signs/traffic_signs_net

```

no-entry.jpg: class=17, No entry



pedestrians.jpg: class=18, General caution



speed-limit-30.jpeg: class=1, Speed limit (30km/h)



stop_sign.jpg: class=14, Stop



warning.png: class=18, General caution



yield.jpg: class=13, Yield



Analyze Performance

```
In [10]: ### Calculate the accuracy for these 5 new images.  
### For example, if the model predicted 1 out of 5 signs correctly, it'  
true_labels=np.array([17, 27, 1, 14, 18, 13])  
  
predictions = predictions  
accuracy = np.mean(predictions==true_labels)  
print("Accuracy: %.2f" % accuracy)
```

Accuracy: 0.83

Output Top 5 Softmax Probabilities For Each Image Found on the

Web

For each of the new images, print out the model's softmax probabilities to show the **certainty** of the model's predictions (limit the output to the top 5 probabilities for each image).

`tf.nn.top_k` (https://www.tensorflow.org/versions/r0.12/api_docs/python/nn.html#top_k) could prove helpful here.

`tf.nn.top_k` will return the values and indices (class ids) of the top k predictions. So if k=3, for each sign, it'll return the 3 largest probabilities (out of a possible 43) and the corresponding class ids.

```
In [11]: ### Print out the top five softmax probabilities for the predictions on  
### Feel free to use as many code cells as needed.  
pd.set_option("display.width", 1000)  
pd.options.display.float_format = '{:,.3g}'.format  
with tf.Session() as sess:  
    probs = tf.nn.softmax(tf.constant(logits))  
    topk_probs, topk_classes = sess.run(tf.nn.top_k(probs, k=5))  
    names = [os.path.basename(f) for f in image_files]  
    df = pd.DataFrame({"name": names})  
    for i in range(5):  
        df["prob_%d" % i] = topk_probs[:,i]  
    df["class_0"] = [signnames[id] for id in topk_classes[:,0]]  
    df["class_1"] = [signnames[id] for id in topk_classes[:,1]]  
    print(df)
```

	name	prob_0	prob_1	prob_2	prob_3	prob_4
class_0	class_1					
0	no-entry.jpg	1	2.49e-16	2.42e-17	8.68e-19	6.78e-22
No entry	Bumpy road					
1	pedestrians.jpg	0.996	0.00408	7.39e-08	4.71e-08	5.28e-09
General caution	Pedestrians					
2	speed-limit-30.jpeg	0.945	0.0552	0.000227	5.27e-08	1.1e-08
Speed limit (30km/h)	Speed limit (80km/h)					
3	stop_sign.jpg	1	0.000401	2.76e-05	5.3e-06	1.01e-06
Stop	Speed limit (80km/h)					
4	warning.png	1	4.94e-07	1.38e-07	7.19e-08	1.37e-09
General caution	Pedestrians					
5	yield.jpg	1	2.24e-23	6.85e-26	2.25e-27	7.29e-28
Yield	Speed limit (30km/h)					

Project Writeup

Once you have completed the code implementation, document your results in a project writeup using this [template \(https://github.com/udacity/CarND-Traffic-Sign-Classifier-Project/blob/master/writeup_template.md\)](https://github.com/udacity/CarND-Traffic-Sign-Classifier-Project/blob/master/writeup_template.md) as a guide. The writeup can be in a markdown or pdf file.

Note: Once you have completed all of the code implementations and

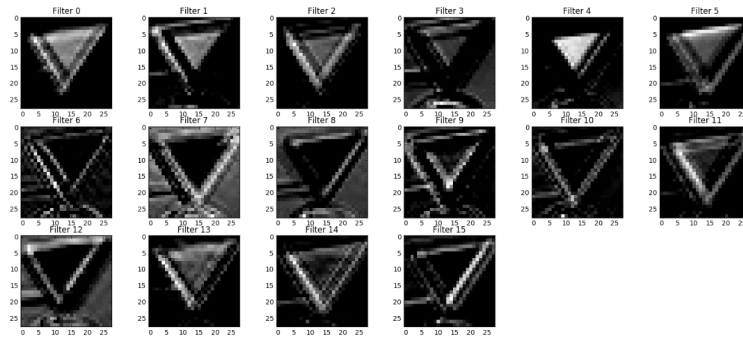
successfully answered each question above, you may finalize your work by exporting the iPython Notebook as an HTML document. You can do this by using the menu above and navigating to "\n", "**File -> Download as -> HTML (.html)**". Include the finished document along with this notebook as your submission.

Step 4 (Optional): Visualize the Neural Network's State with Test Images

This Section is not required to complete but acts as an additional exercise for understanding the output of a neural network's weights. While neural networks can be a great learning device they are often referred to as a black box. We can understand what the weights of a neural network look like better by plotting their feature maps. After successfully training your neural network you can see what its feature maps look like by plotting the output of the network's weight layers in response to a test stimuli image. From these plotted feature maps, it's possible to see what characteristics of an image the network finds interesting. For a sign, maybe the inner network feature maps react with high activation to the sign's boundary outline or to the contrast in the sign's painted symbol.

Provided for you below is the function code that allows you to get the visualization output of any tensorflow weight layer you want. The inputs to the function should be a stimuli image, one used during training or a new one you provided, and then the tensorflow variable name that represents the layer's state during the training process, for instance if you wanted to see what the [LeNet lab's](https://classroom.udacity.com/nanodegrees/nd013/parts/6bf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/601ae704-1035-4287-8b11-e2c2716217ad/concepts/d4aca031-508f-4e0b-b493-e7b706120f81) (<https://classroom.udacity.com/nanodegrees/nd013/parts/6bf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/601ae704-1035-4287-8b11-e2c2716217ad/concepts/d4aca031-508f-4e0b-b493-e7b706120f81>) feature maps looked like for its second convolutional layer you could enter conv2 as the tf_activation variable.

For an example of what feature map outputs look like, check out NVIDIA's results in their paper [End-to-End Deep Learning for Self-Driving Cars](https://devblogs.nvidia.com/parallelforall/deep-learning-self-driving-cars/) (<https://devblogs.nvidia.com/parallelforall/deep-learning-self-driving-cars/>) in the section Visualization of internal CNN State. NVIDIA was able to show that their network's inner weights had high activations to road boundary lines by comparing feature maps from an image with a clear path to one without. Try experimenting with a similar test to show that your trained network's weights are looking for interesting features, whether it's looking at differences in feature maps from images with or without a sign, or even what feature maps look like in a trained network vs a completely untrained one on the same sign image.



Your output should look something like this (above)

```
In [12]: ### Visualize your network's feature maps here.
### Feel free to use as many code cells as needed.

# image_input: the test image being fed into the network to produce the
# tf_activation: should be a tf variable name used during your training
# activation_min/max: can be used to view the activation contrast in mo
# plt_num: used to plot out multiple different weight feature map sets

def outputFeatureMap(image_input, tf_activation, activation_min=-1, act:
    # Here make sure to preprocess your image_input in a way your netwo
    # with size, normalization, ect if needed
    # image_input =
    # Note: x should be the same name as your network's tensorflow data
    # If you get an error tf_activation is not defined it may be having
    activation = tf_activation.eval(session=sess, feed_dict={x : image_i
    featuremaps = activation.shape[3]
    plt.figure(plt_num, figsize=(15,15))
    for featuremap in range(featuremaps):
        plt.subplot(6,8, featuremap+1) # sets the number of feature map
        plt.title('FeatureMap ' + str(featuremap)) # displays the featu
        if activation_min != -1 & activation_max != -1:
            plt.imshow(activation[0,::, featuremap], interpolation="ne
        elif activation_max != -1:
            plt.imshow(activation[0,::, featuremap], interpolation="ne
        elif activation_min != -1:
            plt.imshow(activation[0,::, featuremap], interpolation="ne
        else:
            plt.imshow(activation[0,::, featuremap], interpolation="ne
```